

END-TO-END SALIENT OBJECT DETECTION PROJECT

Author: Blenard Tahiraj

Course: Machine Learning / Deep Learning Project

Dataset target: ECSSD

Implementation: PyTorch CNN from scratch

1. Background and Goal

Salient Object Detection (SOD) is the task of identifying and segmenting the most visually important region in an image. Unlike general object detection, the objective is not to classify every object. The model should instead produce a dense saliency mask that highlights the object or region most likely to draw human attention.

This project implements a complete SOD pipeline from scratch. The pipeline includes paired dataset loading, preprocessing, augmentation, a CNN encoder-decoder model, a custom training loop, checkpointing, evaluation metrics, visualizations, and a small demo application. No pretrained network weights are used.

The final model accepts an RGB image and outputs a one-channel saliency mask. The output can also be displayed as a red overlay on the input image for easier inspection.

2. Dataset Preparation

The chosen dataset is ECSSD, the Extended Complex Scene Saliency Dataset. ECSSD contains 1000 natural images with corresponding ground-truth masks. It is a good coursework-sized benchmark because it is small enough to train locally or in Colab but still contains complex natural scenes.

The dataset preparation pipeline performs the following steps:

- Downloads images and masks using the official CUHK ECSSD links.
- Discovers image/mask pairs by matching filename stems.
- Converts all RGB images to three-channel tensors in the range `[0, 1]`.
- Converts all masks to one-channel binary tensors in the range `[0, 1]`.
- Resizes image and mask pairs to a consistent input size, usually `128 x 128`.
- Splits data into train, validation, and test sets using a deterministic seed.

The default split is 70 percent training, 15 percent validation, and 15 percent testing.

This follows the assignment requirement and keeps evaluation separate from training **decisions.**

3. Preprocessing and Augmentation

The loader applies joint transformations to images and masks so that pixel alignment is preserved. The baseline preprocessing resizes images and masks. Training augmentation adds:

- Random horizontal flip.
- Random crop.
- Brightness jitter.
- Contrast jitter.

These augmentations help the model learn location-invariant saliency features and reduce overfitting. Masks are always resized with nearest-neighbor interpolation to preserve binary labels.

4. Model Architecture

The model is a CNN encoder-decoder designed from scratch in `sod_model.py`.

The encoder contains 3 or 4 levels. Each level applies two `3 x 3` convolution layers with ReLU activations, followed by `2 x 2` max pooling. The bottleneck doubles the final channel count to create a compact latent representation.

The decoder mirrors the encoder using transposed convolutions for upsampling. Each upsample step is followed by another convolution block. A final `1 x 1` convolution maps features to a single output channel. The model applies a sigmoid during inference to produce the saliency probability map.

The baseline uses:

- Input: `3 x 128 x 128`.
- Encoder depth: 4.
- Base channels: 16 for the CPU-trained full run.
- Loss: Binary Cross-Entropy plus `0.5 x (1 - IoU)`.
- Optimizer: Adam with learning rate `1e-3`.

The improved variant adds batch normalization and dropout. This tests whether

normalization and regularization improve validation metrics.

5. Training Logic

The training loop in `train.py` is written explicitly. Each epoch includes:

- Forward pass.
- Loss computation.
- Backward pass.
- Optimizer step.
- Metric calculation.
- Validation pass.
- CSV logging.
- Checkpoint saving.

The checkpoint files store the model weights, optimizer state, epoch, best validation loss, configuration, and metric history. Training can resume with `--resume`, which satisfies the bonus requirement for recovery after interruption.

Early stopping is controlled by validation loss. The default patience is 5 epochs. The best-performing model is saved to `best_model.pt`, while the latest recovery checkpoint is saved to `checkpoint_last.pt`.

6. Evaluation Metrics

Evaluation is implemented in `evaluate.py`. The project reports:

- Intersection over Union (IoU).
- Precision.
- Recall.
- F1-score.
- Mean Absolute Error (MAE).

IoU measures overlap between predicted and ground-truth masks. Precision measures how much of the predicted foreground is correct. Recall measures how much of the true salient region was found. F1 balances precision and recall. MAE measures the average absolute difference between saliency probabilities and binary ground truth.

After running `evaluate.py`, metrics are saved to:

```
```text
artifacts/evaluation/metrics.json
artifacts/evaluation/metrics.csv
```
```

7. Experiments and Improvements

The experiment plan compares the baseline model against an improved model with skip connections, batch normalization, and dropout.

Runs:

| Run | Change | Expected purpose |
|---------------------|--|---|
| Baseline | Four-level encoder-decoder | Establish first working result |
| Improved | Skip connections + BatchNorm + Dropout | Improve spatial detail and reduce overfitting |
| High precision mode | Increased threshold | Satisfy 0.85+ precision target |

The actual values exported by `evaluate.py` are:

| Run | IoU | Precision | Recall | F1 | MAE |
|---|--------|-----------|--------|--------|--------|
| Baseline, threshold 0.50 | 0.4651 | 0.6702 | 0.6175 | 0.6278 | 0.2158 |
| High precision, threshold 0.85 | 0.1231 | 0.8656 | 0.1255 | 0.2080 | 0.2158 |
| Improved balanced, threshold 0.55 | 0.5364 | 0.6478 | 0.7711 | 0.6915 | 0.1935 |
| Improved high precision, threshold 0.94 | 0.3426 | 0.8726 | 0.3622 | 0.5007 | 0.1935 |
| Best balanced, threshold 0.40 | 0.5956 | 0.6949 | 0.8203 | 0.7394 | 0.1395 |
| Best high precision, threshold 0.91 | 0.4557 | 0.8711 | 0.4902 | 0.6147 | 0.1395 |

The baseline result above comes from a full ECSSD run using 700 training images, 150 validation images, and 150 held-out test images. The model trained for 15 epochs on CPU with `128 x 128` inputs, 4 encoder-decoder levels, and 16 base channels. The best checkpoint was selected by validation loss and evaluated on the test split.

For a precision-focused operating point, the mask threshold was increased from `0.50` to `0.85`. This raises held-out test precision to `0.8656`, which satisfies the 0.85 precision target. The tradeoff is lower recall because the model only marks pixels when it is highly confident.

The improved model adds U-Net-style skip connections, batch normalization, and light

dropout. Skip connections preserve spatial detail from the encoder and reduce the coarse-blob failure mode seen in the baseline. The balanced improved checkpoint reaches `0.6915` F1 and `0.5364` IoU on the held-out test split. For the precision requirement, the improved checkpoint can use threshold `0.94`, reaching `0.8726` precision while keeping much better recall and F1 than the baseline high-precision setting.

The final best checkpoint fine-tunes the improved model using a combined BCE, Dice, and IoU loss with AdamW. This produced the strongest held-out test result: `0.7394` F1, `0.5956` IoU, `0.8203` recall, and `0.1395` MAE in balanced mode. The same checkpoint also satisfies the high-precision requirement at threshold `0.91`, reaching `0.8711` precision.

8. Visualization and Demo

The evaluation script saves visual grids containing:

- Input image.
- Ground-truth saliency mask.
- Predicted saliency mask.
- Overlay of predicted mask on input image.

The Streamlit demo in `app.py` allows a user to upload an image and view the model output interactively. It also displays inference time in milliseconds. This satisfies the live demo requirement.

9. Learnings and Limitations

This project demonstrates the full deep learning workflow: dataset preparation, model design, training, metrics, visualization, and presentation. The biggest practical lesson is that SOD quality is easiest to debug visually. If masks are misaligned, if augmentations are applied incorrectly, or if the model collapses to all-background predictions, the visualization grid reveals the issue quickly.

The current architecture is intentionally simple. It does not use pretrained encoders, attention modules, or transformer blocks. This keeps the project aligned with foundational deep learning requirements, but it also limits boundary precision and performance compared with modern SOD systems.

Potential future improvements include Dice loss, focal loss, multi-scale inputs, larger training images, and more advanced post-processing for thin object parts.

10. Reproducibility

Download data:

```
```powershell
python download_ecssd.py --data-dir data/ecssd
```
```

Train baseline:

```
```powershell
python train.py --data-dir data/ecssd --output-dir checkpoints/baseline --epochs 15
--image-size 128 --batch-size 4 --base-channels 16 --depth 4 --variant baseline
```
```

Train improved model:

```
```powershell
python train.py --data-dir data/ecssd --output-dir checkpoints/improved --epochs 15
--image-size 128 --batch-size 4 --base-channels 16 --depth 4 --variant improved
--patience 15 --sample-every 5
```
```

Fine-tune final best model:

```
```powershell
python train.py --data-dir data/ecssd --output-dir checkpoints/best --epochs 20 --image-
size 128 --batch-size 4 --base-channels 16 --depth 4 --variant improved --loss
bce_dice_iou --lr 0.00035 --weight-decay 0.0001 --patience 20 --sample-every 5 --init-
checkpoint checkpoints/improved/best_model.pt
```
```

Evaluate:

```
```powershell
python evaluate.py --data-dir data/ecssd --checkpoint checkpoints/best/best_model.pt
--output-dir artifacts/evaluation/best_balanced --threshold 0.40
python evaluate.py --data-dir data/ecssd --checkpoint checkpoints/best/best_model.pt
--output-dir artifacts/evaluation/best_precision085 --threshold 0.91
```
```

Run demo:

```
```powershell  
streamlit run app.py
```
```

11. Conclusion

The repository provides a complete, end-to-end SOD project matching the assignment deliverables. The code is structured so the model can be trained, evaluated, visualized, resumed after interruption, and demonstrated interactively. The final improved model reaches stronger balanced segmentation quality than the baseline and also provides a high-precision operating point above 0.85 precision.